

Mutation Readiness Scorecard

Version 1.0-draft · 18 questions · 6 dimensions · ~15 minutes
Source: Appendix A of Mutation Readiness: An Operating Manual for Innovation in the Age of AI by Mark Saymen

Machine-readable source of truth: [mutation-readiness-scorecard.yaml](#). The web assessment, the Claude Code skill, and the MCP server all derive from that file. If you change a question, change it there first.

What this measures

The organization’s capability to **sense, respond, and adapt at AI-age speed**, across six dimensions, three questions each. The unit-of-analysis matters: a 50-person team can be Mutation-Ready while the surrounding 5,000-person enterprise is Mutation-Blind. Decide what you are scoring before you start, and hold every answer to that unit.

How to score

Rate each question 1 (strongly disagree) to 5 (strongly agree). N/A is allowed — reduce that dimension’s denominator proportionally.

- **Dimension score:** sum of the three answers (max 15), converted to the 0-12 headline number by $\text{floor}(\text{sum} \times 12 / 15)$.
- **Total score:** sum of the six dimension scores, out of 72.
- If some questions are N/A, normalize the total to a percentage.

Band	Score	Reading
Mutation-Blind	0-30	Operating on lagging metrics in an environment that is mutating around you. Disruption is probably already

Band	Score	Reading
		happening, undetected.
Mutation-Aware	31-50	You can sense change but cannot yet act on it at the right cadence. The sense-to-respond gap is the primary risk.
Mutation-Ready	51-72	Top decile. The risk now is complacency.

The questions

Dimension 1 — Signal Sensitivity

The capability to detect, interpret, and act on weak signals before lagging metrics confirm them.

- **Q1.1** Does your team monitor weak signals from at least three sources outside your own dashboards on a weekly basis? (Customer Discord channels, developer forums, GitHub trends, competitor hiring patterns, regulatory leaks.)
- **Q1.2** Can you name a specific behavioural change in your customer base that you noticed in the past 30 days before it appeared in your metrics?
- **Q1.3** Do you have a structured cadence — a meeting, a ritual, a tool — for surfacing internal anomalies without political consequence to the person who raised them?

Dimension 2 — Structural Flexibility

The organization's ability to reshape itself faster than competitors can retool.

- **Q2.1** Are your teams aligned to customer outcomes (streams of value) rather than to internal functions?
- **Q2.2** Can a new product idea move from concept to validated learning in under 90 days?
- **Q2.3** When you identified a need to restructure something in the past 12 months, did you act on it within a quarter?

Dimension 3 — AI Talent Flywheel

The capability to attract, retain, and integrate AI-literate talent across functions.

- **Q3.1** Do you have at least one AI-literate person embedded in every major product or business team?
- **Q3.2** Is your AI talent distributed across the organization, or isolated in a single AI/ML function?
- **Q3.3** Does your AI talent have direct exposure to strategy decisions, not just to implementation work?

Dimension 4 — Ambidextrous Capital

The discipline of balancing exploit (proven operations) with explore (uncertain bets) in your funding model.

- **Q4.1** Do you have a protected exploration budget that is separate from your core P&L?
- **Q4.2** Are exploration bets evaluated on learning yield, not on traditional ROI?
- **Q4.3** Can an experiment with no clear ROI survive its first quarterly business review without being defunded?

Dimension 5 — Ethical Guardrails

Containment-as-velocity: the institutional discipline that lets you ship AI fast without dramatic reputational failure.

- **Q5.1** Is every AI deployment governed by explicit operating boundaries and a recalibration cadence?
- **Q5.2** Do you have an AI governance function (not just a compliance function) with engineering rather than legal at its centre?
- **Q5.3** Can you produce, on request, an audit-quality explanation of any AI-driven decision affecting a customer?

Dimension 6 — Narrative Coherence

The shared story that lets the organization act in concert under uncertainty.

- **Q6.1** Can your employees describe, in their own words, what your organization is for in the current market (not the 2018 market)?
- **Q6.2** Has your leadership team updated the strategic narrative in the past 18 months?
- **Q6.3** Can new hires articulate your strategic narrative after their first month on the job?

Rapid mode

Six questions — the strongest single signal per dimension: Q1.2, Q2.2, Q3.2, Q4.3, Q5.3, Q6.1. Score out of 30. Bands: 0-12 Mutation-Blind, 13-22 Mutation-Aware, 23-30 Mutation-Ready.

After scoring

Each band carries recommended actions and a reading list — see the [YAML spec](#) or the [Framework Guide](#). For a guided conversational session with a personalized 30-60-90-day report, use the [Claude Code skill](#).

To score programmatically (18 answers in Q1.1...Q6.3 order; na allowed):

```
python3 score.py --answers 3,4,2,3,3,4,2,2,3,4,3,3,2,2,3,3,4,3
python3 score.py --answers 3,4,na,3,3,4,2,2,3,4,3,3,2,2,3,3,4,3
python3 score.py --rapid --answers 4,3,2,3,4,2
```

Companion instrument

Run the [Innovation Framework Scorecard](#) first if you haven't measured your innovation culture: it scores the cultural operating system underneath these practices. An organization in Innovation Theatre will fail at Mutation Readiness regardless of what it installs.