

The Mutation Readiness Framework Guide

The Mutation Readiness Framework Guide

Version 1.0-draft · June 2026 Status: DRAFT — pending owner review and sign-off before publication.

The definitive guide to the Mutation Readiness Framework: the rules, definitions, instruments, and cadences for building an organization that adapts faster than its market changes.

This guide is licensed under Creative Commons Attribution-ShareAlike 4.0 (CC BY-SA 4.0), pending owner confirmation. You may share and adapt it with attribution, under the same license.

1. Purpose and Use

Every organization is built for the environment that existed when it was designed. Environments no longer hold still long enough for that design to pay back. The arrival of capable AI has compressed the cycle further: the cost of producing software, content, analysis, and coordination is collapsing, and with it the moats built on those costs.

Mutation readiness is an organization's capacity to sense, decide, and reconfigure faster than its environment changes. It is not a transformation program with an end date. It is an operating property, like solvency or safety — something you maintain, measure, and can lose.

This guide is the canonical definition of the framework. Like other reference guides of its kind, it is deliberately minimal: it defines the concepts, the instruments, and the cadence, and leaves the application to practitioners. The companion books — The

Innovation Playground (a novel) and Mutation Readiness: An Operating Manual for Innovation in the Age of AI — supply narrative, evidence, and extended practice. The framework repository supplies the machine-readable instruments. Where any of these disagree with this guide, this guide prevails.

Who this is for. Leaders accountable for an organization's continued relevance; practitioners running the framework inside teams; facilitators teaching it; and builders integrating it into tools (the entire framework is published in machine-readable form for that purpose).

How to use it. Read it once end to end (it is short by design). Then run the Mutation Readiness Scorecard, plot your portfolio on the 4×4 Innovation Matrix, install the weekly signal review, and reassess quarterly. That loop — assess, plot, review, reassess — is the whole operating system. Everything else is technique.

2. Definitions

Mutation. A change in the environment significant enough that the organization's current design is no longer the right answer. Mutations are not projects, trends, or quarters; they are shifts in what works. The emergence of capable generative AI is a mutation. So was the smartphone. So is a regulatory regime change in your industry.

Mutation readiness. The organization's capability to sense, respond, and adapt at AI-age speed — measured across six dimensions (§3) on a 0-72 scale with three bands:

- **Mutation-Blind (0-30):** the organization is operating on lagging metrics in an environment that is mutating around it. Disruption is probably already happening, undetected.
- **Mutation-Aware (31-50):** the organization can sense change but cannot yet act on it at the right cadence. The sense-to-respond gap is the primary risk.
- **Mutation-Ready (51-72):** the organization is operating in the top decile. The risk now is complacency.

Signal. A discrete, observable piece of evidence that a mutation may be under way: a customer behaving differently, a technology crossing a capability threshold, a competitor making an otherwise-irrational move, a regulation drafted. Signals are logged, not debated, at the point of detection (see the SignalNet template).

Weak signal. A signal supported by a single source or speculative interpretation. Weak signals are the cheapest information an organization will ever acquire about its future; the framework's first rule is that logging them must cost almost nothing.

Leading indicator. A measure that moves before an outcome does, while intervention is still possible. Every dimension and every initiative in the framework carries at least one leading indicator (see the Leading Indicator Dashboard template). Lagging indicators are kept for accountability; leading indicators are kept for steering.

Experiment. A time-boxed, cost-capped initiative with kill criteria written before funding. An initiative without kill criteria is a commitment, which is legitimate — but the framework requires the organization to say which one it is making.

Kill criteria. The evidence, defined in advance with a date, that ends an experiment. Kill criteria written after a pitch has been approved are theater.

3. The Six Dimensions

Mutation readiness is measured across six dimensions. Each is necessary; none is sufficient. The dimensions are deliberately cross-functional — no single department can own mutation readiness, and an organization that assigns it to one has already failed the assessment.

1. **Signal Sensitivity.** The capability to detect, interpret, and act on weak signals before lagging metrics confirm them. Operationalized by the SignalNet — a distributed, semi-anonymous log any employee can write to — and a weekly, anomaly-focused signal review.
2. **Structural Flexibility.** The organization's ability to reshape itself faster than competitors can retool: teams aligned to streams of value rather than functions, concept-to-validated-learning in under 90 days, restructuring decisions acted on within a quarter.
3. **AI Talent Flywheel.** The capability to attract, retain, and integrate AI-literate talent across functions — embedded in every major team rather than isolated in a single AI/ML silo, and exposed to strategy decisions, not just implementation work.

4. **Ambidextrous Capital.** The discipline of balancing exploit (proven operations) with explore (uncertain bets) in the funding model: a protected exploration budget separate from the core P&L, bets evaluated on learning yield, experiments that can survive their first quarterly business review.
5. **Ethical Guardrails.** Containment-as-velocity: the institutional discipline that lets you ship AI fast without dramatic reputational failure. Explicit operating boundaries and a recalibration cadence for every AI deployment; a governance function with engineering rather than legal at its centre; audit-quality explanations on request.
6. **Narrative Coherence.** The shared story that lets the organization act in concert under uncertainty: employees who can say what the company is for in the current market, a strategic narrative refreshed within the past 18 months, new hires who can articulate it after their first month.

The full instrument — 18 questions, three per dimension, with scoring rubric — is defined in §6 and published in machine-readable form in the framework repository (`assessments/mutation-readiness-scorecard.yaml`).

4. The 4×4 Innovation Matrix

The matrix plots every initiative in the portfolio on two axes.

Ambition (columns): Optimize — make the existing thing measurably better; Extend — take it to new users, markets, or uses; Transform — replace how it works while keeping what it is for; Mutate — change what the organization is for, obsoleting part of yourself before the market does.

Unit of change (rows): Product — what you sell; Process — how work gets done; Business model — how value is captured; Organization — how people, structure, and incentives are arranged.

The matrix exists to make portfolio shape visible. Three rules govern its use:

1. **Every active initiative gets exactly one cell.** Arguments about which cell are not friction; they are the diagnosis surfacing.
2. **Healthy portfolios hold positions in at least three of the four ambition columns at all times.** A portfolio entirely in Optimize/Extend is optimizing itself into irrelevance; a portfolio entirely in Transform/Mutate has mistaken adrenaline for strategy.

3. **The Mutate column is funded only with capped, kill-criteria experiments** — never from the core P&L, and ideally shielded structurally (separate team, separate rules, separate building if necessary).

Initiatives drift left over time as they mature. The drift is normal; unnoticed drift is how transformation programs quietly become maintenance programs. Re-plot quarterly.

5. The Five Levers of Reinvention

Draft note: the canonical Five Levers are defined in Chapter 9 of Mutation Readiness: An Operating Manual for Innovation in the Age of AI. The lever names below are a working draft and must be reconciled against Chapter 9 by the author before v1.0 final.

When an assessment reveals a gap, these are the five levers an organization can actually pull. Every intervention the framework prescribes is a setting of one or more of these levers; if a proposed action does not map to a lever, it is commentary, not intervention.

1. **Capital.** Move money on a rolling cadence instead of an annual one. Reallocation ratio — the percentage of innovation budget that moved this quarter — is the lever's gauge. Zero for three consecutive quarters means the lever is rusted shut.
2. **Cadence.** Install the operating rhythm: weekly signal review, quarterly reassessment and portfolio re-plot, annual recalibration (§7). Cadence is the cheapest lever and the most commonly skipped, because calendars are where intentions go to die.
3. **Capability.** Build AI fluency and experimentation skill deliberately: embedded in real workflows with measured before/after deltas, not training catalogs. Capability investments are experiments too — they carry leading indicators and review dates.
4. **Configuration.** Make structure cheap to change: modular processes, value-stream funding, teams that can form and dissolve in weeks. The gauge is the last actual reconfiguration, timed end to end.
5. **Conviction.** Leadership behavior and incentives. Leaders run their own experiments and publish their own position changes; compensation rewards adaptation. This lever cannot be delegated, which is why it is listed last and matters first.

6. The Mutation Readiness Scorecard

The scorecard is the framework's core instrument: **18 questions, three per dimension, each rated 1-5** (strongly disagree → strongly agree), N/A allowed. It is published in full in Appendix A of the Operating Manual and in machine-readable form in the framework repository.

Scoring rubric. Dimension score = sum of the three answered questions (max 15), converted to the 0-12 headline number by $\text{floor}(\text{sum} \times 12 / 15)$; with N/A items the denominator reduces proportionally. Total = sum of the six dimension scores, out of **72**, reported alongside a six-axis profile (radar chart in rendered reports); if any items are N/A the total is also normalized to a percentage. Band thresholds: Mutation-Blind 0-30, Mutation-Aware 31-50, Mutation-Ready 51-72.

Administration. Choose the unit of analysis first — yourself, a team, or the whole organization — and hold every answer to it. A 50-person team can be Mutation-Ready inside a Mutation-Blind enterprise, and a “yes” that applies only to a sub-team is a 3, not a 5. Three modes, in increasing order of value:

- Personal scan (~15 minutes): one leader answers. Fast, free, and systematically optimistic — useful as a door, not a diagnosis.
- Team scan: each member of a leadership team answers independently before seeing each other's responses. The spread between answers is itself a primary finding; a team whose scores on Narrative Coherence diverge by more than 4 points has located its real agenda.
- Quarterly trend: the same team scan repeated on the recalibration cadence. The trend line matters more than any absolute score.

A rapid mode exists for screening: six questions, one per dimension (Q1.2, Q2.2, Q3.2, Q4.3, Q5.3, Q6.1), scored out of 30 with bands 0-12 / 13-22 / 23-30. It is a screen, not a diagnosis.

Interpretation rules.

- The lowest dimension, not the total, sets the work program. Dimensions compound; a 12-12-12-12-12-2 profile underperforms a 9-9-9-9-9-9 one.
- Each band carries prescribed actions and a reading list (published in the scorecard specification), delivered as a 30-60-90-day plan in facilitated reports. A key Mutation-Aware metric: **mutation latency**, the time from signal identification to organizational decision — target under 21 days.
- Scores are self-reported and will inflate under audience pressure. Never attach compensation or public league tables

to scorecard results; the instrument measures honesty as much as readiness, and incentivized scores measure neither.

The full question text, band definitions, and actions are normative and published in the repository at `assessments/mutation-readiness-scorecard.{md,yaml}`. The companion **Innovation Framework Scorecard** (15 questions across the five Principles of the Innovators Framework: Fail Forward, Open Communication & Candid Feedback, Challenging Mental Models, Flexible Structure / Agile Decision-Making, Growth Mindset at Scale; scored /75 into Innovation Theatre / Practicing / Native) measures the cultural operating system underneath these practices. Run it first: an organization in Innovation Theatre will fail at mutation readiness regardless of what it installs.

7. Recommended Cadence

The framework runs on three loops. All three are required; the framework does not function as an annual event.

Weekly — signal review (30 minutes). Fixed agenda: new signals (10 minutes, 60 seconds each, no debate), status changes on watched signals (10 minutes), and at most one escalation decision (10 minutes). One owner, one log, same slot every week. This meeting is the framework's heartbeat; if only one practice survives contact with your calendar, keep this one.

Quarterly — reassessment and re-plot (half day). Re-run the scorecard as a team scan; re-plot the portfolio on the 4×4 matrix; review the leading indicator dashboard and the KPI scorecard; audit the quarter's surprises against the SignalNet (was the miss a sensing failure or a deciding failure?); set or reset one lever (\$5).

Annual — recalibration (one to two days). Step back from the loops: re-examine whether the dimensions are weighted right for your context, retire gamed indicators, review every kill decision of the year for false positives, and decide deliberately what the organization will stop sensing — attention is the scarcest budget the framework spends.

8. Glossary

Term	Definition
Band	One of three scored ranges on the scorecard: Mutation-Blind (0-30), Mutation-Aware (31-50), Mutation-Ready (51-72).

Term	Definition
Dimension	One of the six measured components of mutation readiness (§3).
Experiment	A time-boxed, cost-capped initiative with pre-written kill criteria.
Kill criteria	Evidence, defined before funding, that ends an experiment by a set date.
Kill latency	How long an initiative survives past its own kill criteria. Target: zero.
Leading indicator	A measure that moves before the outcome does, while action is still possible.
Lever	One of the Five Levers of reinvention (Operating Manual, Chapter 9; working draft names in §5).
Mutation	An environmental change large enough that the organization's current design is no longer the right answer.
Mutation exposure	The cost of inaction if a signal proves real; the systematically underpriced side of the risk ledger.
Mutation latency	Time from signal identification to organizational decision. Target: under 21 days.
Reallocation ratio	Percentage of innovation budget that moved between initiatives in a quarter.
Signal	A discrete, observable piece of evidence that a mutation may be under way.
SignalNet	A distributed, semi-anonymous logging system for any employee to surface weak signals, with a weekly review cadence (Operating Manual, Appendix E).
Team scan	Scorecard mode where each leader answers independently; the spread is a primary finding.

9. Acknowledgments and References

The Mutation Readiness Framework was developed by Mark Saymen and is elaborated in The Innovation Playground and Mutation Readiness: An Operating Manual for Innovation in the Age of AI.

The framework integrates work by Amy Edmondson (psychological safety, intelligent failure), Mary Murphy (cultures of growth), Carol Dweck (mindset), Kim Scott (radical candor), Ethan Mollick (co-intelligence), Mustafa Suleyman (containment), Marco Iansiti and Karim Lakhani (the AI factory), Rita McGrath (inflection points), Adam Grant (hidden potential), Matthew Skelton and Manuel Pais (Team Topologies), Peter Senge (the learning organization), and the Anthropic interpretability research from 2024–2025.

It also stands on ground prepared by the lean startup and evidence-based innovation tradition (build-measure-learn, innovation accounting); portfolio approaches to innovation ambition (McKinsey's three horizons and their critics); the agile and lean-portfolio body of practice, including the Scaled Agile Framework's treatment of lean budgeting and value streams; the organizational-learning literature on psychological safety (Edmondson) and dynamic capabilities (Teece); and the format example of the Scrum Guide, which demonstrated that a framework's canonical text can and should fit in a short, versioned document.

Contributions, translations, and case studies are welcome — see CONTRIBUTING.md in the framework repository.

Changelog - v1.0-draft (June 2026) — initial draft for owner review. Dimensions, questions, scoring rubric, bands, and band actions reconciled with the canonical assessment (Appendix A of the Operating Manual). Five Lever names (§5) and 4×4 matrix cell archetypes (§4) remain working drafts requiring owner sign-off before v1.0 final.